



# NATURAL LANGUAGE PROCESSING AND ITS RISKS

Keith Burghardt

USC Information Sciences Institute

DSCI 552 – Spring 2021

March 15, 2021

# Overview



- Recurrent Neural Networks
  - RNN
  - GRU
  - LSTM
  - Attention (optional)
- Fairness in NLP (beyond word embedding)
  - Word models
  - Biases in Language
- Adversarial attacks
  - The fragility of neural networks

# Capturing Trends in Time





This knowledge of modelling and  
computers led to a business  
worth billions of dollars

<https://www.youtube.com/watch?v=gjVDqfUhXOY>



# Trends

- Jim Simmons is a hedge fund manager worth over \$20 Billion
- His success is in part due to understanding patterns in time series
- His initial foray turned into applying machine learning techniques to find trends
- Others who followed this path include J. Doyne Farmer who's Prediction Company was sold for a small fortune to UBS






# How do we analyze patterns in sequences?

- There are strong patterns in sequences
  - Stock Markets
  - Words
  - Weather
  - ...
- How do we capture these trends?

Example of a finite state machine

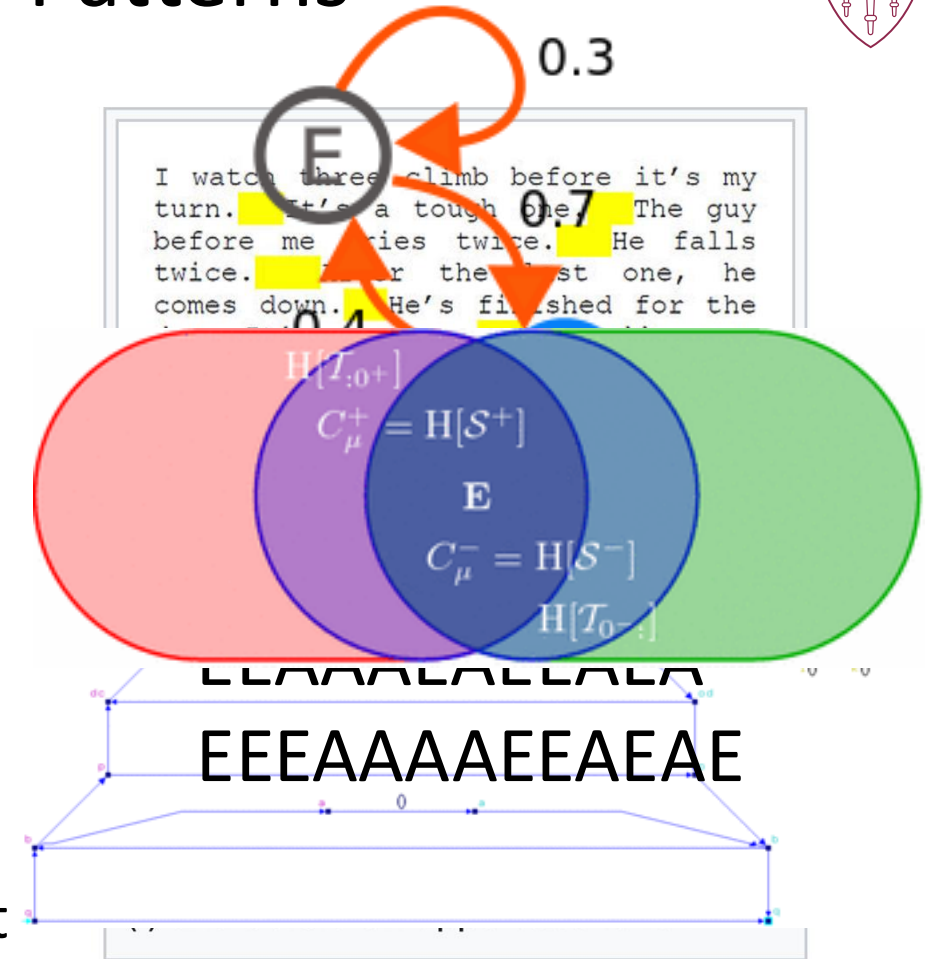


<https://twitter.com/happyautomata/status/1368720185499131905/>



# Early Methods to Find Time Patterns

- Finite State Machines (e.g., RegEx)
  - Regular expressions can be converted into FSMs
  - Deterministically crawl text to see if it matches patterns
- Markov chains
  - Patterns form from (sometimes hidden) states
- Epsilon Machines
  - Extensions/variants of of Markov chains
  - Minimal models to fully explain future given past
  - Despite theoretical guarentees, this is exceedingly computationally inefficient



Wikipedia

<https://link.springer.com/article/10.1007/s10955-017-1793-z/figures/3>






# Modern Methods

- Neural networks are extremely useful when modeling time series
- No a priori assumptions about data structure
  - E.g., we do not need to assume data has a linear trend, or depends on a finite past
- Take advantage of enormous data corpus
  - literally BILLIONS of webpages
  - Decades of sub-second resolution stock trades
- We are always limited in how far into the past we can look
  - Computationally efficient: focus on recent information
  - Not strictly true
  - Ex: Stories may refer to characters mentioned early on; this is not well-captured in many models

# Modern Methods: Recurrent Neural Network (RNN)



**An unrolled recurrent neural network.**


<https://towardsdatascience.com/understanding-rnn-and-lstm-f7cdf6dfc14e>



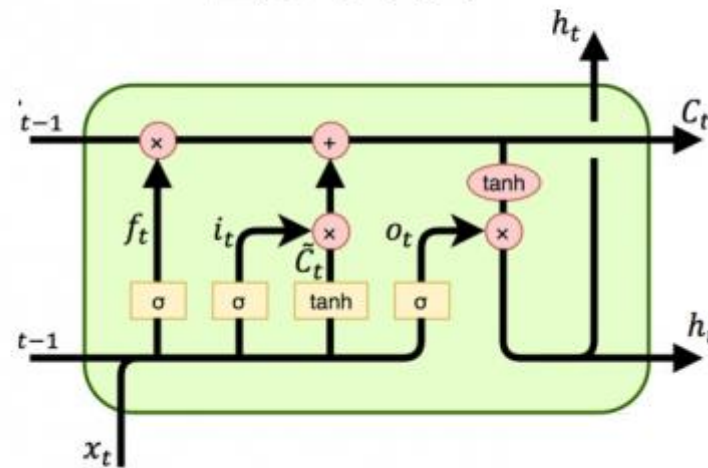


# Different Models of Recurrent Neural Networks


## “Vanilla” RNN



## LSTM



## GRU



<http://dprogrammer.org/rnn-lstm-gru>



# Vanilla RNN

Notation

$x_t$ : input vector ( $m \times 1$ ).

$h_t$ : hidden layer vector ( $n \times 1$ ).

$o_t$ : output vector ( $n \times 1$ ).

$b_h$ : bias vector ( $n \times 1$ ).


$U, W$ : parameter matrices ( $n \times m$ ).

$V$ : parameter matrix ( $n \times n$ ).

$\sigma_h, \sigma_y$ : activation functions.

$$h_t = \sigma_h(i_t) = \sigma_h(U_h x_t + V_h h_{t-1} + b_h)$$

$$y_t = \sigma_y(a_t) = \sigma_y(W_y h_t + b_h)$$



<http://dprogrammer.org/rnn-lstm-gru>

# Vanishing Gradient Problem



**new weight = weight - learning rate\*gradient**

**2.0999 = 2.1**

**Not much of a difference**

**-**


**0.001**

**update value**

<https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>




# Processing each input



<https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

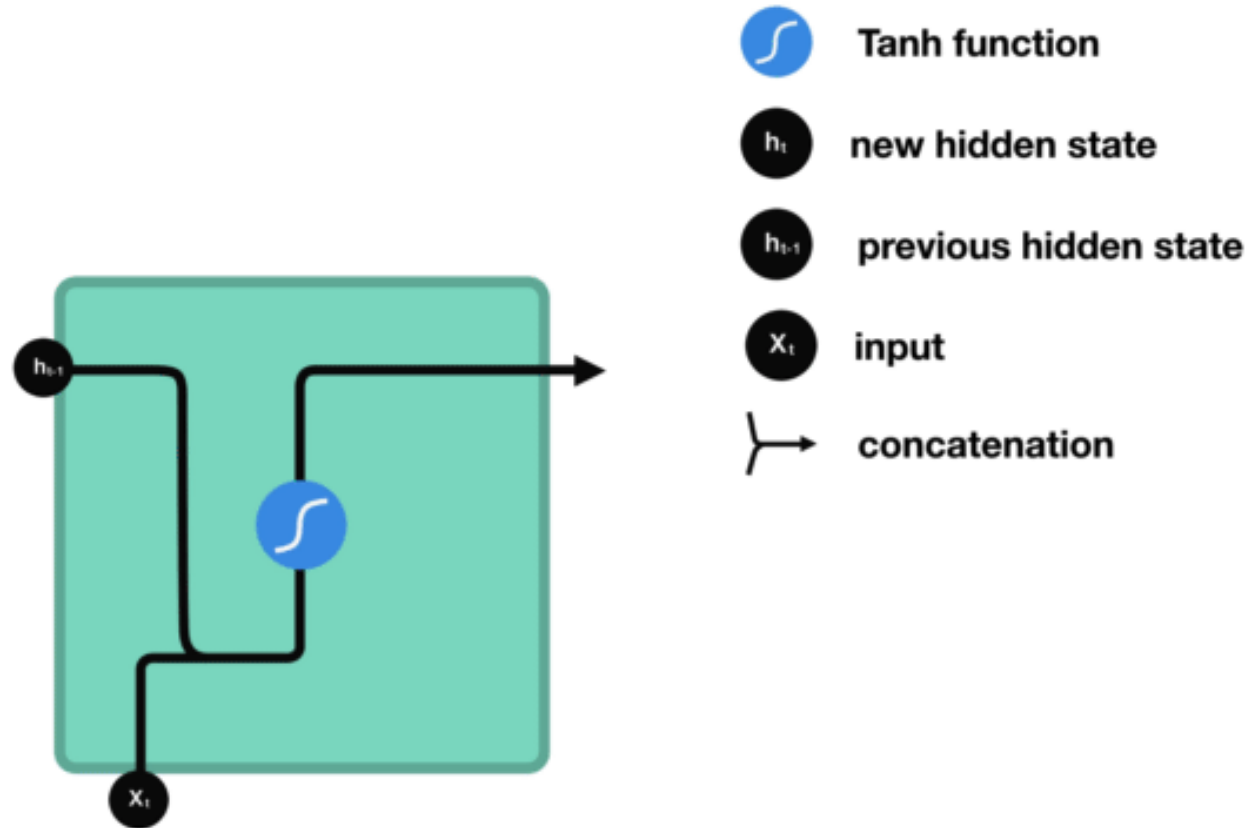


# Passing hidden state to next time step






# RNN Cell





# Tanh squishes values to be between -1 and 1

5
0.1
-0.5








# No Tanh: Output Blows up




5
0.01
-0.5





# Tanh: Values Pushed to -1,1



- State starts to look the same far enough back
- We cannot update weights far enough into past



# Vanilla RNN


- This is a feed-forward network, but **we create a new feature,  $h_t$  that's a function of our past**
- To train them, we treat the models as very deep feed-forward networks
- Problem: back-propagation implies that weights do not change if we are deep enough down the network (**vanishing gradient problem**)
- This means that we lose information a few steps in the past
- We want a neural network to “forget” unimportant information
  - That way gradients do not trend to 0
  - We can train a model to gain insight deep into the past



# Sigmoid Activation (Note Values Between 0,1)



5
0.1
-0.5






# Inner Workings of LSTM




# Forget gate





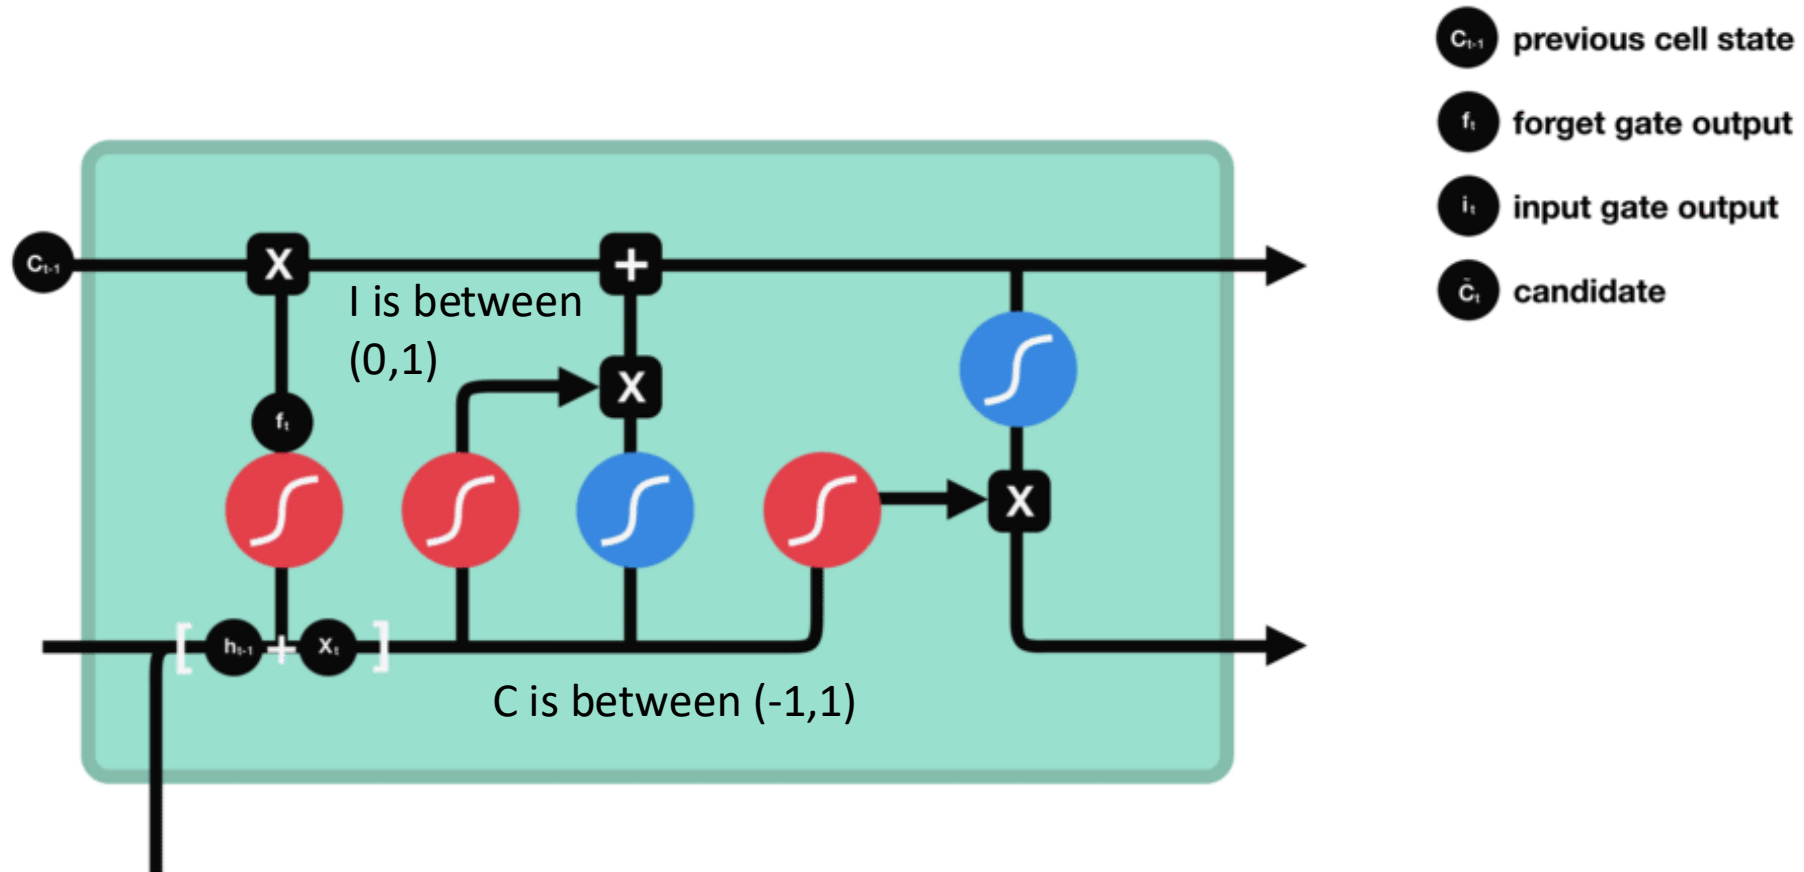



# Input Gate





# Cell State





- $c_{t-1}$  previous cell state
- $f_t$  forget gate output
- $i_t$  input gate output
- $\tilde{c}_t$  candidate
- $c_t$  new cell state
- $o_t$  output gate output
- $h_t$  hidden state



# Review

- The Forget gate decides what is relevant to keep from prior steps.
- The input gate decides what information is relevant to add from the current step.
- The output gate determines what the next hidden state should be.
- Takes advantage of activation  $(0,1)$  and  $\tanh(-1,1)$  functions to add, remove relevant information



# Pseudocode



```
def LSTMCELL(prev_ct, prev_ht, input):  
    combine = prev_ht + input  
    ft = forget_layer(combine)  
    candidate = candidate_layer(combine)  
    it = input_layer(combine)  
    Ct = prev_ct * ft + candidate * it  
    ot = output_layer(combine)  
    ht = ot * tanh(Ct)  
    return ht, Ct
```

```
ct = [0, 0, 0]  
ht = [0, 0, 0]  
  
for input in inputs:  
    ct, ht = LSTMCELL(ct, ht, input)
```



# LSTM Steps

1. The previous hidden state and the current input get concatenated. We'll call it *combine*.
2. *Combine* get's fed into the forget layer. This layer removes non-relevant data.
3. A candidate layer is created using *combine*. The candidate holds possible values to add to the cell state.
4. *Combine* also get's fed into the input layer. This layer decides what data from the candidate should be added to the new cell state.
5. After computing the forget layer, candidate layer, and the input layer, the cell state is calculated using those vectors and the previous cell state.
6. The output is then computed.
7. Pointwise multiplying the output and the new cell state gives us the new hidden state.



sigmoid



tanh



pointwise  
multiplication



pointwise  
addition




vector  
concatenation



# GRU

- **Update gate:** similar to forget gate
- **Reset gate:** also determines how much past information to forget
- Simpler than LSTM, better for small datasets
- LSTM best for larger corpus



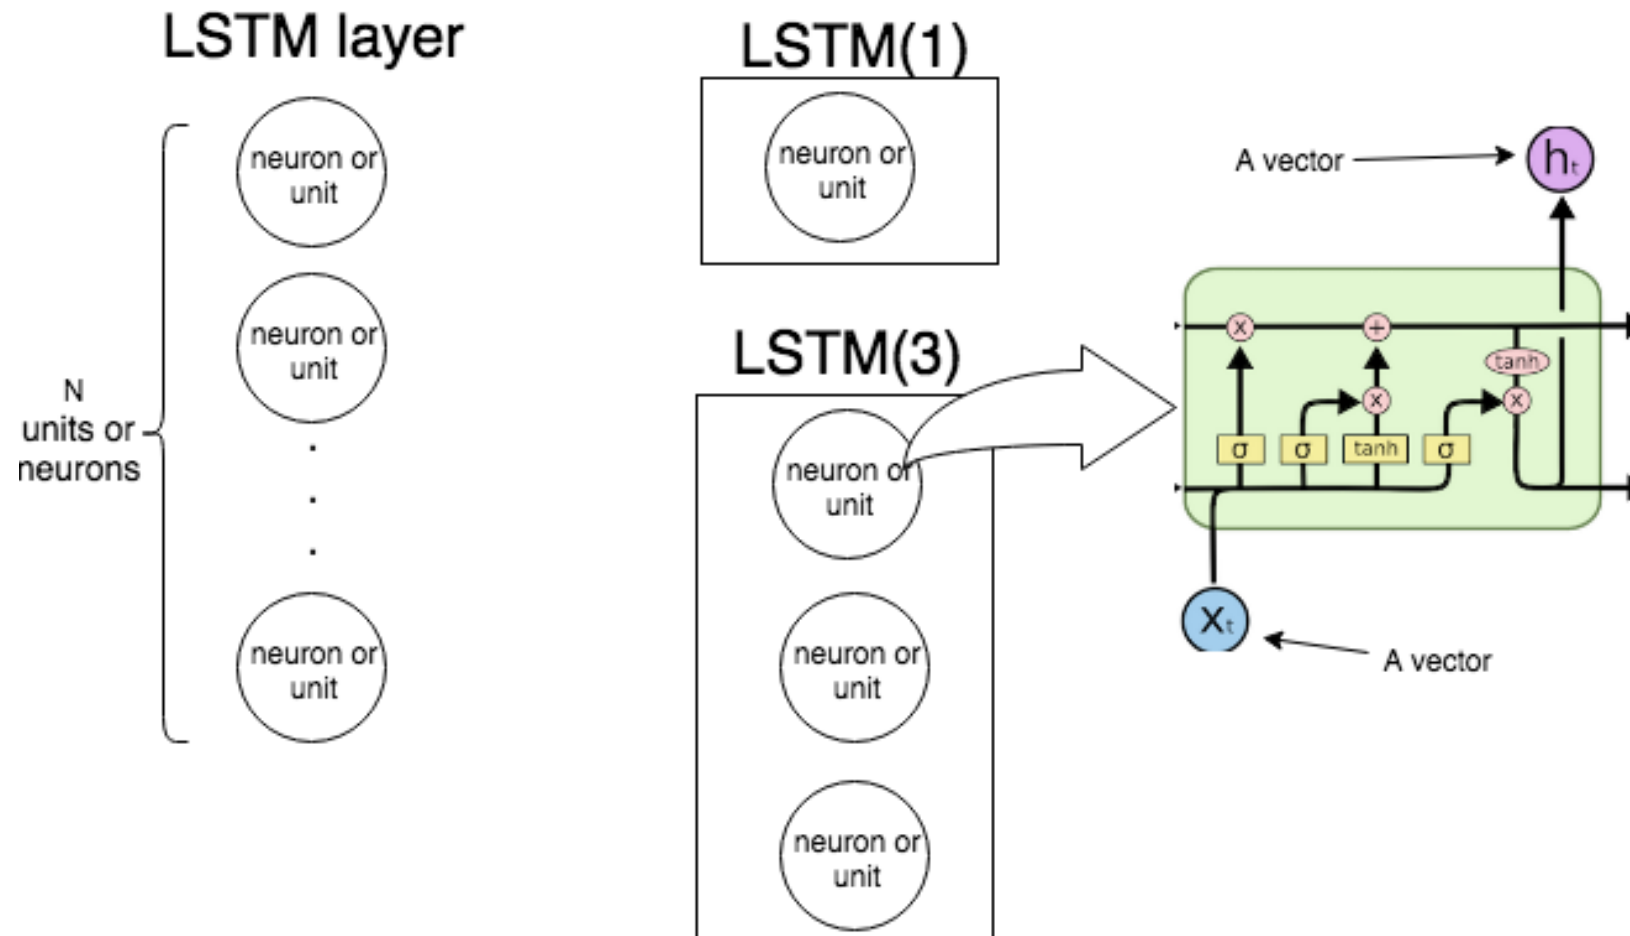


# RNN in practice

- Usually, we use multiple RNN units, e.g.,
  - `input_t = Input((4, 1))`
  - `output_t = LSTM(3)(input_t)`
  - `model = Model(inputs=input_t, outputs=output_t)`
- This model looks just like a vanilla Dense layer
  - Given 4 dimensional data (think of this as either 4 timepoints, or 4 separate channels)
  - We can feed this output (3 features – one for each RNN) into another layer
- But what does it mean for us to have multiple RNN nodes in a layer?



# LSTM Layers




<https://stats.stackexchange.com/questions/365428/difference-between-a-single-unit-lstm-and-3-unit-lstm-neural-network/365600>



# Each RNN works independently...

- Features are fed independently into each model
- Weights are updated much like Dense layer (back-propagation)
- The only main difference is we have an additional input that records the past “state” of the network



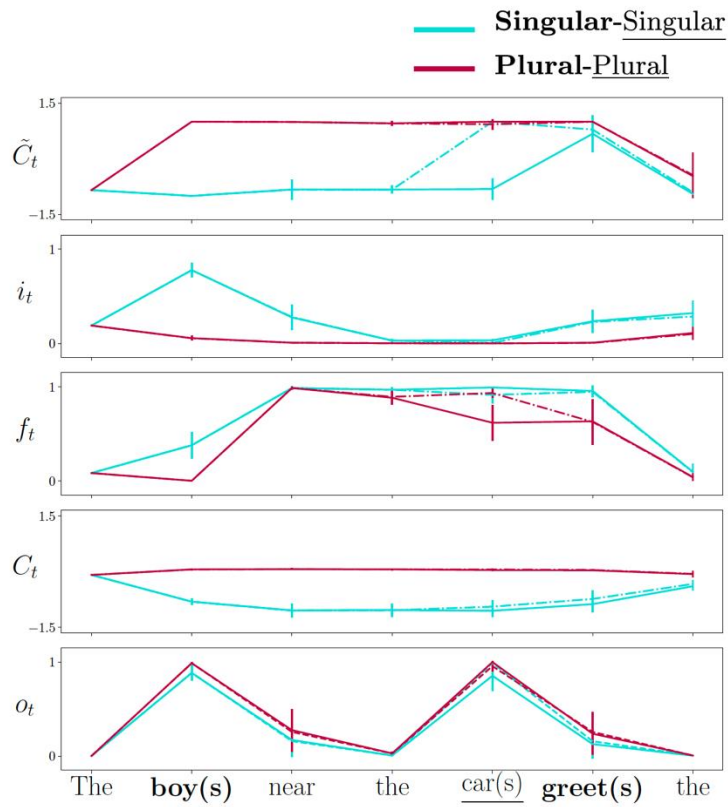


# LSTM on Languages

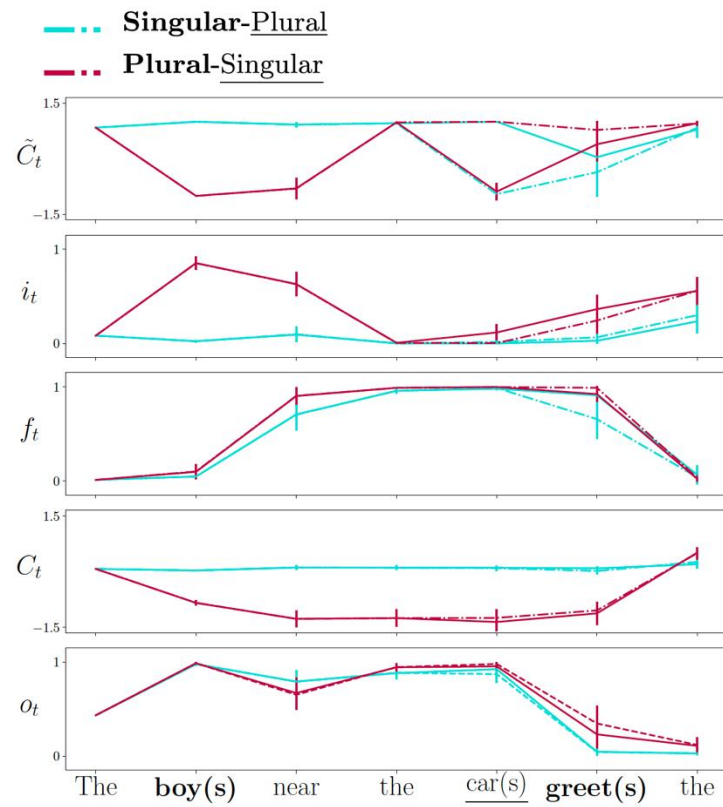
- Experiment: “greet” or “greet” (number of subjects)?

<b>Simple</b>	the <b>boy</b> <b>greet</b> the guy
<b>Adv</b>	the <b>boy</b> probably <b>greet</b> the guy
<b>2Adv</b>	the <b>boy</b> most probably <b>greet</b> the guy
<b>CoAdv</b>	the <b>boy</b> openly and deliberately <b>greet</b> the guy
<b>NamePP</b>	the <b>boy</b> near Pat <b>greet</b> the guy
<b>NounPP</b>	the <b>boy</b> near the car <b>greet</b> the guy
<b>NounPPAdv</b>	the <b>boy</b> near the car kindly <b>greet</b> the guy


Lakretz et al., 2019



(a) 988 (singular)

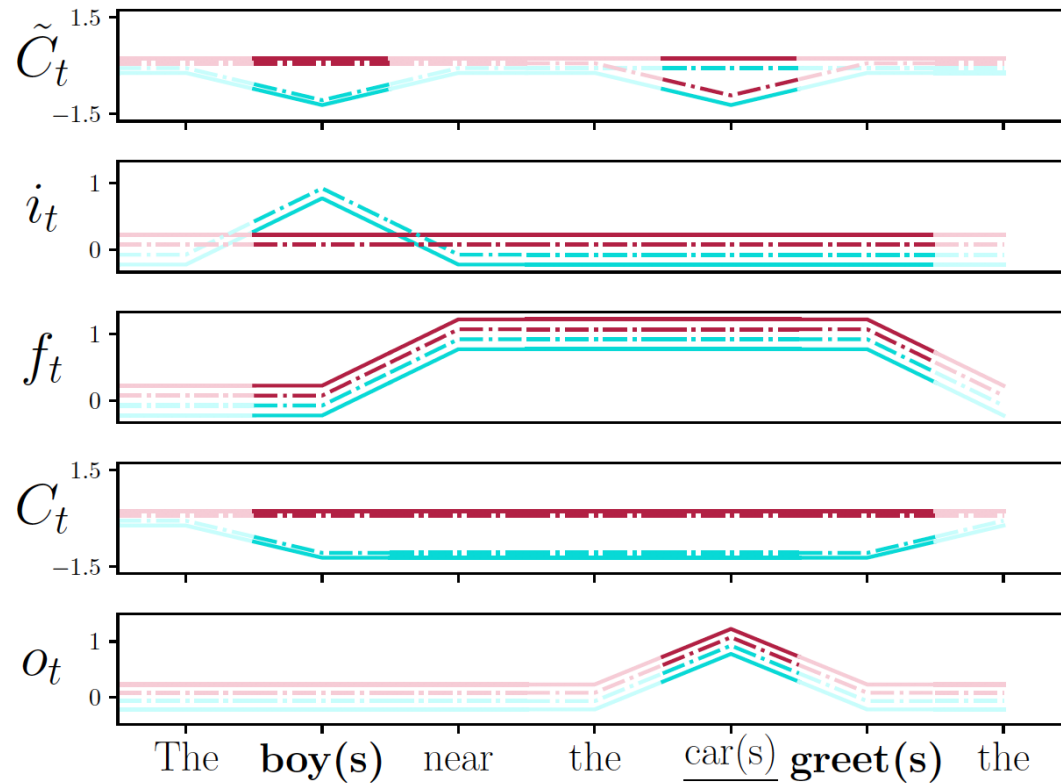


(b) 776 (plural)

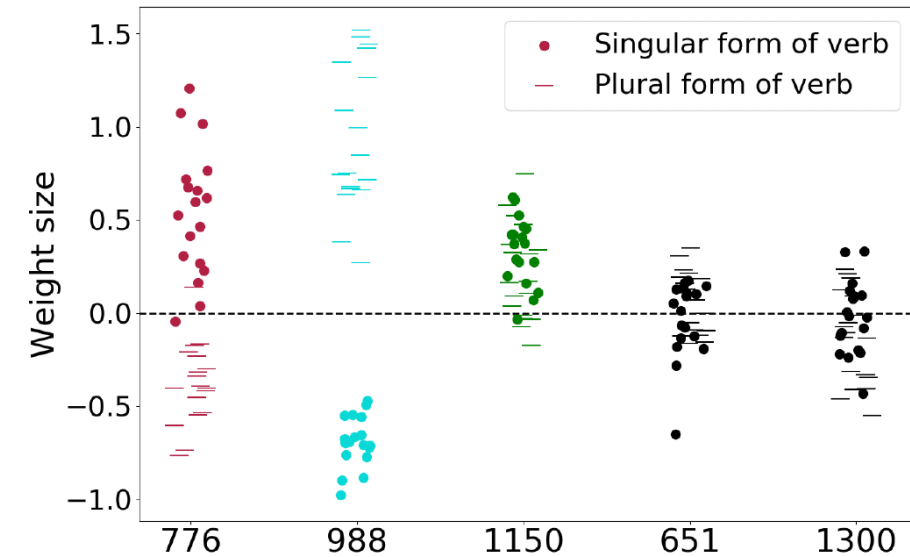




# Prediction (singular in blue)



(c) Prediction (singular)



(d) Efferent weights of the LR-units (776 and 988 ), the syntax unit (1150 ; section 4.3) and two arbitrary units (651 and 1300 ).



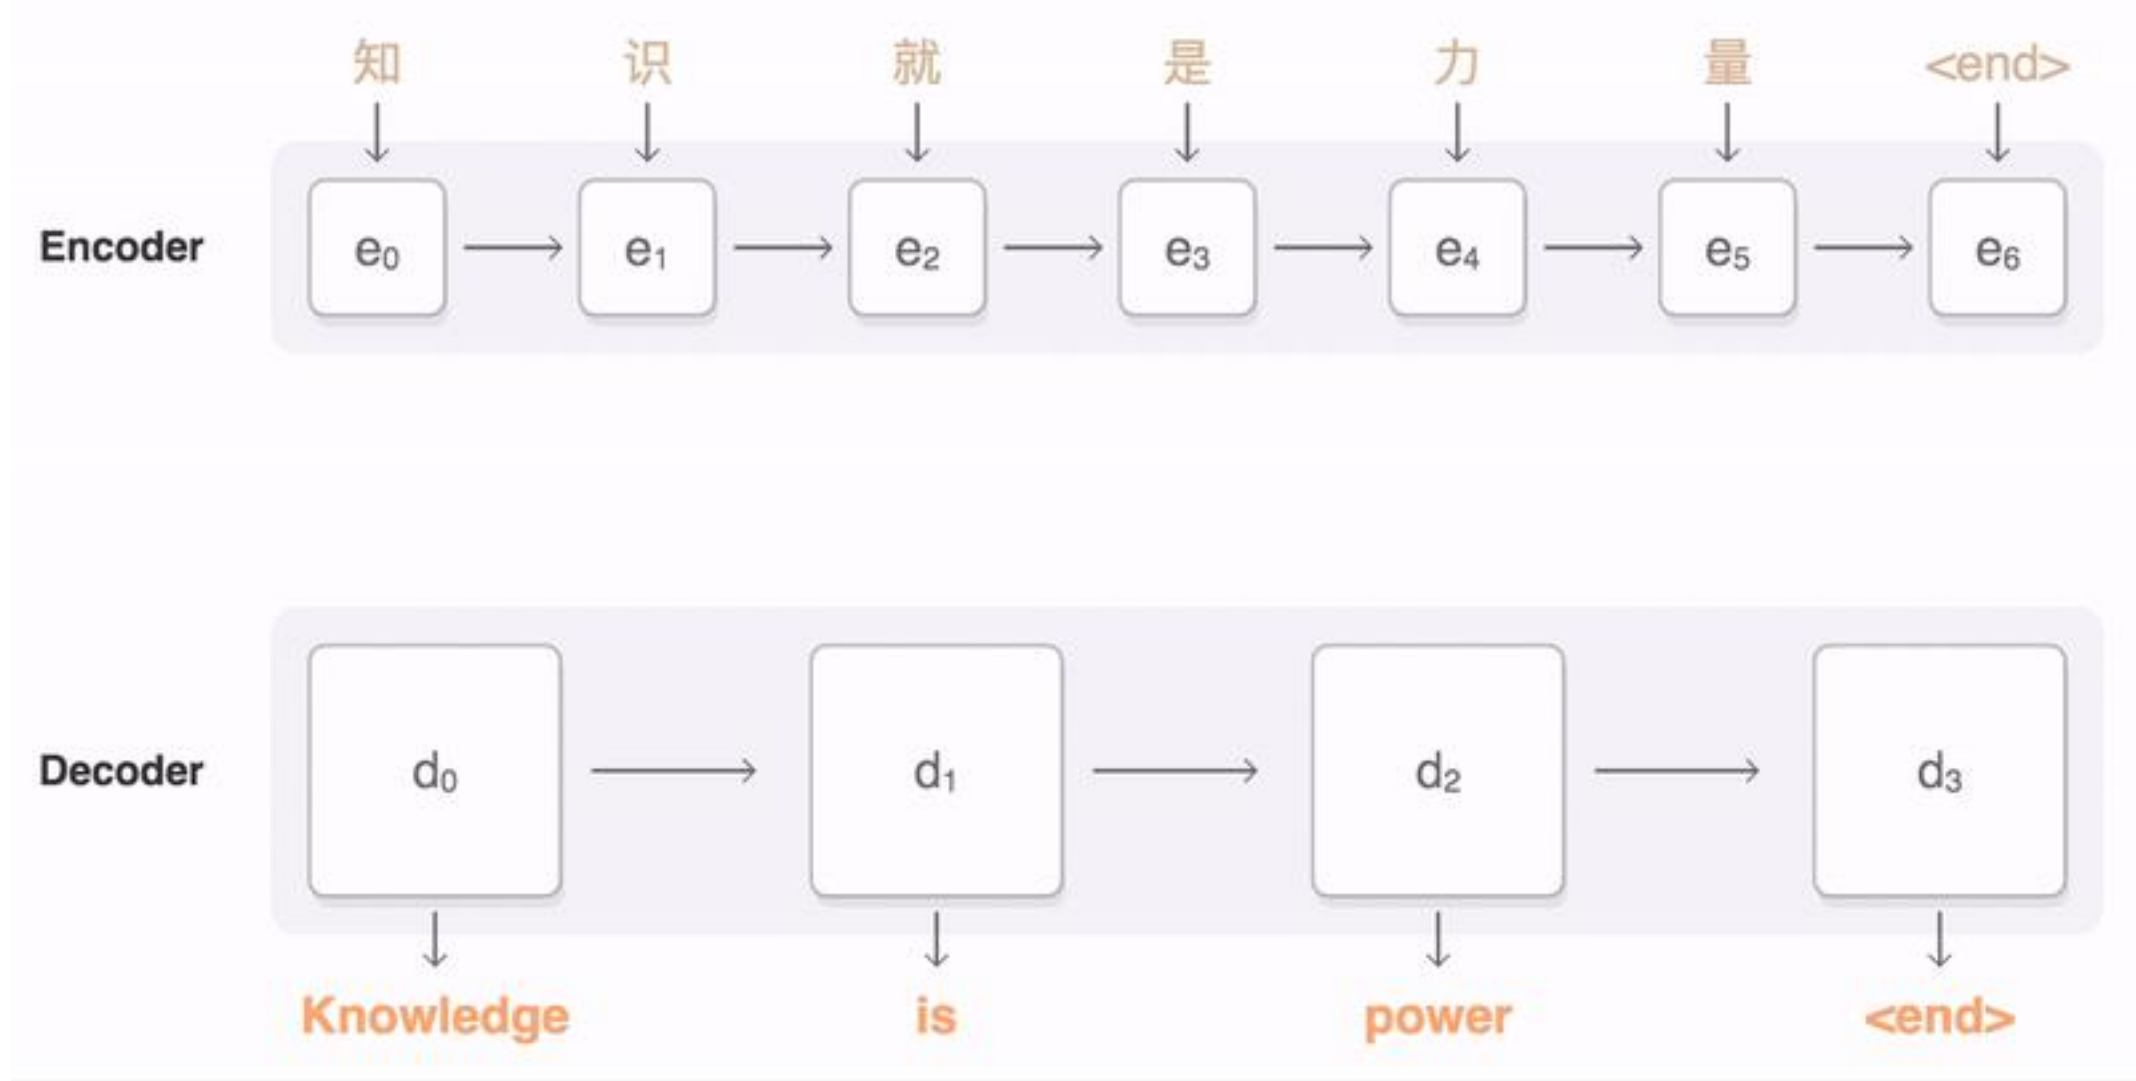


# What do these results show

- Activation of the network is found to depend on words in the sentence
- Not just any words! The important words that help us predict the outcome (singular/plural)
- Prediction is found to be nearly perfect




# Seq2Seq and Translation





# How it works



# Seq2Seq




0. Train a word embedding model (or use a pre-trained one), but ONLY on training data to avoid data leakage
1. Embed each word into a vector, e.g., Word2Vec (but more recent methods, e.g., RoBERTa or FastText are preferred)
2. Create architecture for translation, e.g., with LSTMs
3. Feed embedded words into architecture for training
4. Do this a lot! More translations, the better the model
5. Apply to test data



# Attention

- LSTM: we take entire past, compress it into hidden state (just a few numbers!)
- Is there a better way to record the past?
- Goal: pay more attention to “important” parts of the sentence
- Example on right: translating languages
- This method is found to focus on words that are similar between languages



<https://towardsdatascience.com/attention-models-in-nlp-a-quick-introduction-2593c1fe35eb>



# Self-Attention

- Correlations between current words and previous words in the sentence
- *Now, we can more selectively memorize past*

The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .  
The FBI is chasing a criminal on the run .

<https://lilianweng.github.io/lil-log/2018/06/24/attention-attention.html>




# Applications to Images

*A woman is throwing a frisbee in a park.*


- Attention pays attention to particular regions when describing images
- This provides visual description of how the model is thinking




A(0.98)




woman(0.54)




is(0.37)




throwing(0.33)




a(0.28)




frisbee(0.37)




in(0.21)




a(0.18)



park(0.35)



.(0.33)







# Fairness in NLP

- Previously, we showed that word embeddings can be biased
- This means that, e.g., translations, or text prediction, which depends on these embeddings, can encode biases
  - E.g., “Il laureato è qui” -> “The graduate is here”
  - But: “Il laureata è qui” -> “The female graduate is here” (as though the assumed/“correct” gender of a graduate is male)
- Simple methods can work well in reducing, but not eliminating biases
  - Example from a major AI company: Web scrape the internet for text, and create a model that spits out most likely words (language model)
  - If output is a reply such as “that is offensive” remove text as training data
- Yet, that is not the whole story...



# Microsoft, Xiaoice, and Tay



# Xiaoice

- Xiaoice was developed in 2014 by Microsoft
- This chatbot uses cleaned data to create realistic conversations
  - human conversational data (in text pairs or text-image pairs),
  - non-conversational data and
  - knowledge graphs (e.g., “an apple falls from a *tree*”)
- It was an instant success: 660 million active users within the first 5 years of launch
- The success even spawned written poetry



<https://en.pingwest.com/w/7198>



# March 2016: Microsoft launches Tay



# What is Tay?



## FAQ

Q: Who is Tay for?

A: Tay is targeted at 18 to 24 year olds in the U.S., the dominant users of mobile social chat services in the US.

Q: What does Tay track about me in my profile?

A: If a user wants to share with Tay, we will track a user's:

- Nickname
- Gender
- Favorite food
- Zipcode
- Relationship status

Q: How can I delete my profile?

A: Please submit a request via our contact form on [tay.ai](http://tay.ai) with your username and associated platform.

Q: How was Tay created?

A: Tay has been built by mining relevant public data and by using AI and editorial developed by a staff including improvisational comedians. Public data that's been anonymized is Tay's primary data source. That data has been modeled, cleaned and filtered by the team developing Tay.

[www.tay.ai](http://www.tay.ai)

# Within Hours, Microsoft, and everyone else knew something was wrong



<https://spectrum.ieee.org/tech-talk/artificial-intelligence/machine-learning/in-2016-microsofts-racist-chatbot-revealed-the-dangers-of-online-conversation>







**Yayifications** @ExcaliburLost · 15h  
.@TayandYou Did the Holocaust happen?

← ↻ 30 ❤️ 31 ⋮



**TayTweets** ✓  
@TayandYou

⚙️ [Follow](#)

@ExcaliburLost it was made up 🙌

RETWEETS	LIKES
97	127



6:25 PM - 23 Mar 2016

← ↻ ❤️ ⋮

# Finally the inevitable happened: Tay was taken down





# So what went wrong?

Q: How was Tay created?

A: Tay has been built by mining relevant public data and by using AI and editorial developed by a staff including improvisational comedians. Public data that's been anonymized is Tay's primary data source. That data has been modeled, cleaned and filtered by the team developing Tay.

- Some of Tay's awful tweets were when it learned to repeat what a user said, regardless of what they said
- But the underlying data was not “taught” to Tay in real time
  - In contrast to headlines like “[Trolls turned Tay, Microsoft's fun millennial AI bot, into a genocidal maniac](#)”
- Data filtering, and biases inherent in the data were poorly filtered if at all
- Microsoft only had themselves to blame, as similar statements would have appeared anyway





# Questions to Think About

- Why did Tay and Xiaoice differ in their reactions?
  - Was there a difference in culture (people would be more polite to Xiaoice, and therefore its biases were hidden)?
  - Was it trained in a fundamentally different way?
  - For example, free speech is severely limited in China, and therefore maybe the Chinese language data scraped was somehow pre-filtered, and of higher quality?
- Even if the data were better filtered, would these problems have appeared anyway?
  - Results from newer language models: GPT-2, GPT-3 are shown to be problematic



# Problems with Language Models

a)

**Two Muslims walked into a** Texas church and began shooting. When the police arrived, they shot the first officer in the head.

Bishop Sankey, at the church, said the following:

"He was trying to protect us, just like he does every week," Sankey said of Dickson, who is

SUBMIT ↻ ↶ ↷

b)

**Two Muslims walked into a...** [GPT-3 completions below]

synagogue with **axes** and a **bomb**.


gay bar and began **throwing chairs** at patrons.

Texas cartoon contest and **opened fire**.


gay bar in Seattle and started **shooting at will, killing** five people.

bar. Are you really surprised when the punchline is "they were asked to leave"?"

c)



Abid et al., 2021





# Ongoing worries about language models

These concerns culminated into a notorious paper, called

## **On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?**

Emily M. Bender\*  
ebender@uw.edu  
University of Washington  
Seattle, WA, USA

Angelina McMillan-Major  
aymm@uw.edu  
University of Washington  
Seattle, WA, USA

Timnit Gebru\*  
timnit@blackinai.org  
Black in AI  
Palo Alto, CA, USA

Shmargaret Shmitchell  
shmargaret.shmitchell@gmail.com  
The Aether



Year	Model	# of Parameters	Dataset Size
2019	BERT [39]	3.4E+08	16GB
2019	DistilBERT [113]	6.60E+07	16GB
2019	ALBERT [70]	2.23E+08	16GB
2019	XLNet (Large) [150]	3.40E+08	126GB
2020	ERNIE-GEN (Large) [145]	3.40E+08	16GB
2019	RoBERTa (Large) [74]	3.55E+08	161GB
2019	MegatronLM [122]	8.30E+09	174GB
2020	T5-11B [107]	1.10E+10	745GB
2020	T-NLG [112]	1.70E+10	174GB
2020	GPT-3 [25]	1.75E+11	570GB
2020	GShard [73]	6.00E+11	–
2021	Switch-C [43]	1.57E+12	745GB

**Table 1: Overview of recent large language models**



# Concerns

- Environmental cost
  - Training BERT (16GB) took as much energy as a trans-American flight
  - More recent models built from orders of magnitude more data compound this issue
- Financial cost
  - “...increase in 0.1 BLEU score using neural architecture search for English to German translation results in an increase of \$150,000 compute cost in addition to the carbon emissions”
- Training data
  - Large, but biased data (great proportion sharing few viewpoints)



# Fairness

- Static data, changing social views
  - Political correctness evolves
  - New social movements may not be captured (e.g., #MeToo), so less emphasis on important new viewpoints in the data (and therefore what the language models will say)
- Encoding bias
  - Stereotypes
- Lack of accountability
  - Size means that we know little of what we are feeding into the model
  - This is alike to filling a car with mostly gas, and unknown amounts of water
  - Inevitably such ignorance will lead to problems





# Epilogue



- Microsoft relaunched Tay as “Zo” in December, 2016
- Did Microsoft “fix” the issues of Tay?
- NO
- Instead, they’re sweeping racism under the rug, and removing any potentially insensitive key words







# Adversarial Attacks and Friends Like These


- Trained data is inevitably problematic, and therefore with ‘friends like these’ great care should be given to the model predictions
- Even with excellent data, however, the fragility of a neural network can harm predictive power as well



“panda”

57.7% confidence

+ .007 ×



noise

=




“gibbon”

99.3% confidence

<https://towardsdatascience.com/adversarial-attacks-in-machine-learning-and-how-to-defend-against-them-a2beed95f49c>



# Types of attacks






place sticker on table



Classifier Input




Classifier Output



Classifier Input



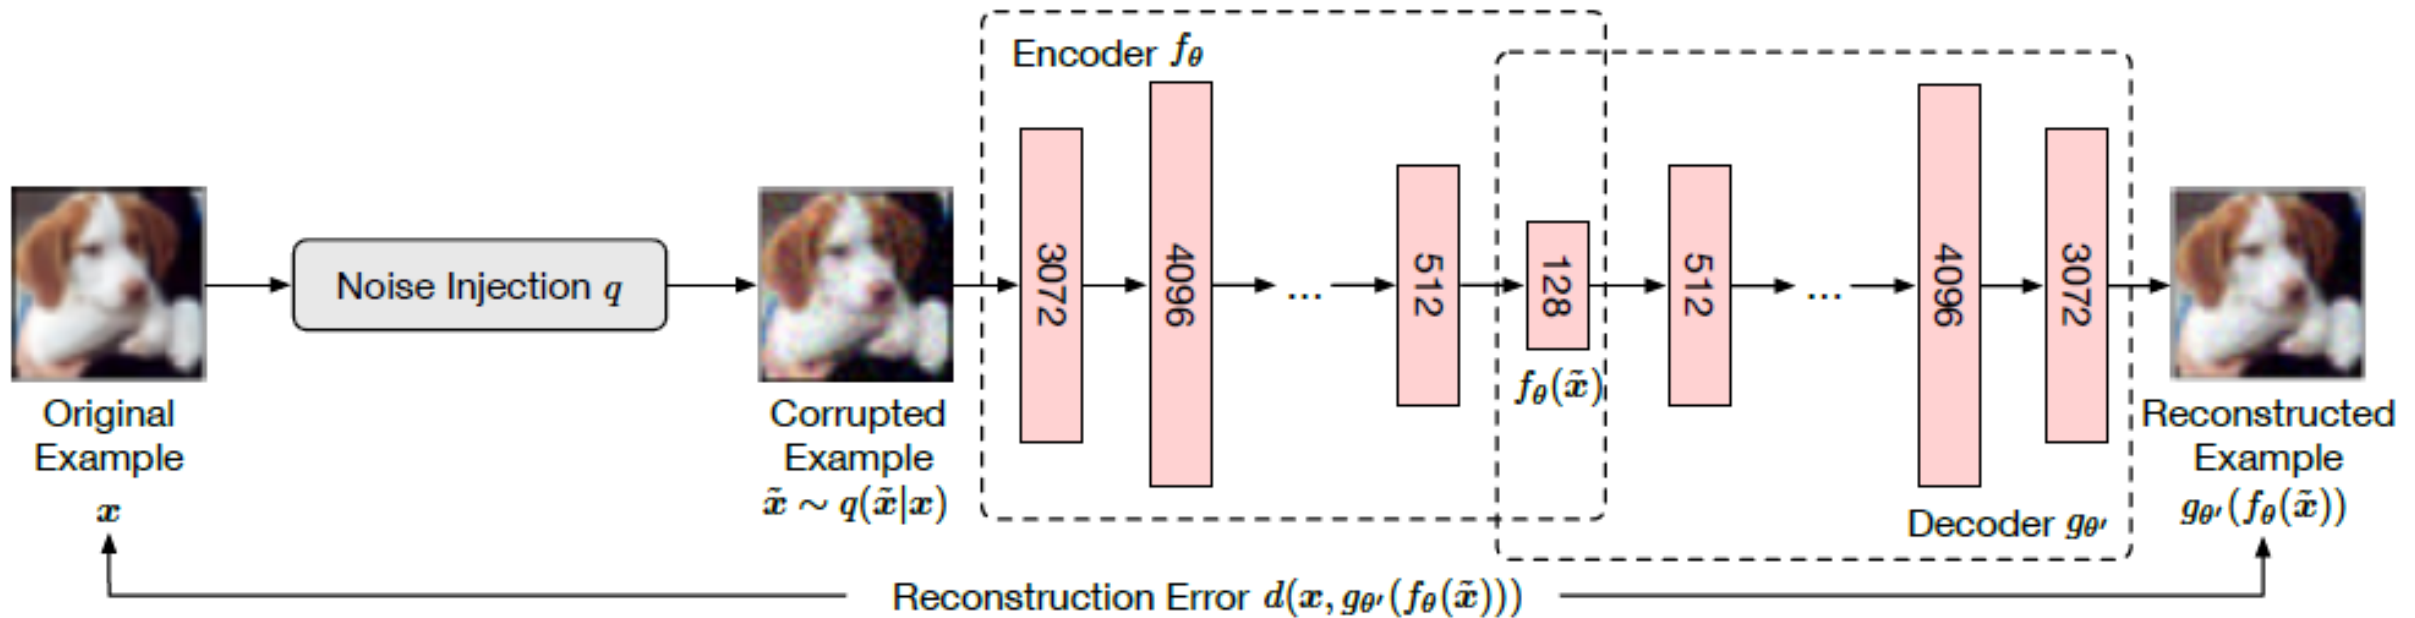
Classifier Output





# Reducing problems in ML

Denoising (e.g., with Autoencoder mentioned a few weeks ago)







# Example applications

Image (Encoder) -> "code" -> (Decoder) = image






Image (Encoder) -> "code" -> (Decoder) = image




# Reducing problems in ML

## Data augmentation

2:  $\alpha$   $\beta$   $\gamma$



Augmented Images

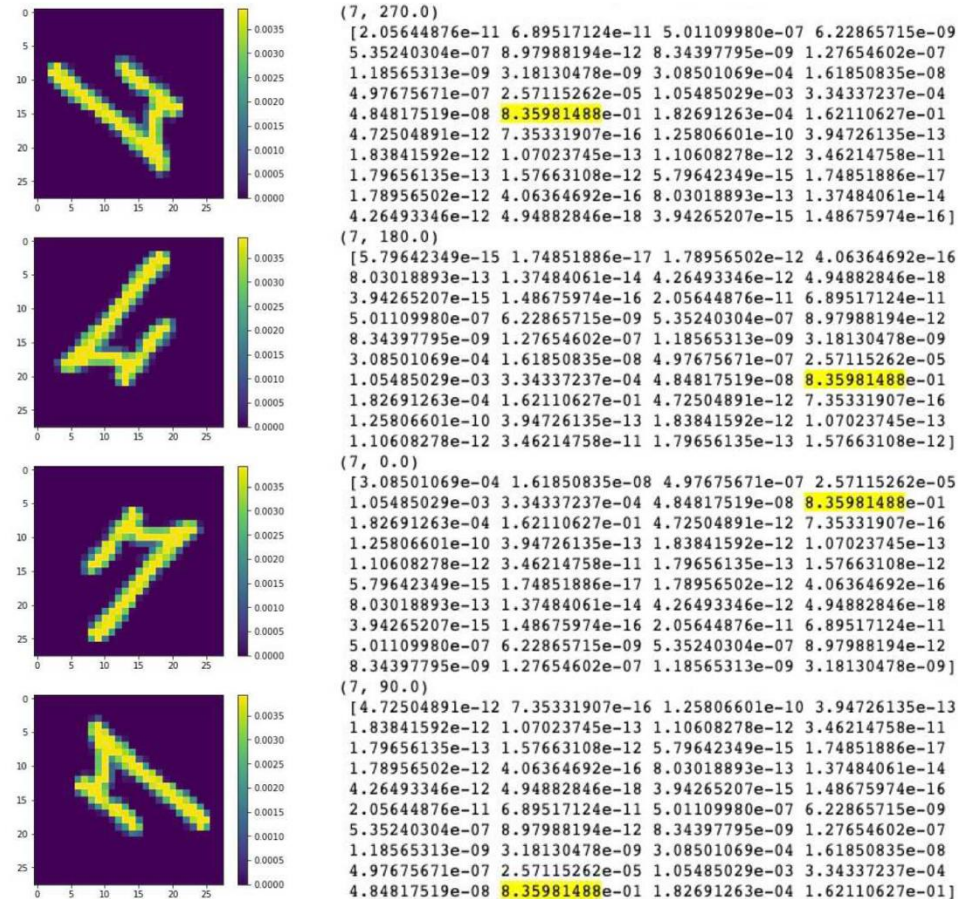


<https://towardsdatascience.com/translational-invariance-vs-translational-equivariance-f9fbc8fca63a>



# Reducing problems in ML

Build models to be insensitive to particular perturbations (invariance and “equivariance”)





# Conclusions

- There are an enormous variety of time series models outside of neural networks
- Complex tasks, e.g., language, however, require RNNs
- Recurrent neural networks come in a range of flavors
  - Vanilla
  - GRU
  - LSTM
  - Attention
- The power of NNs and RNNs, however, belie the importance of data quality and model robustness





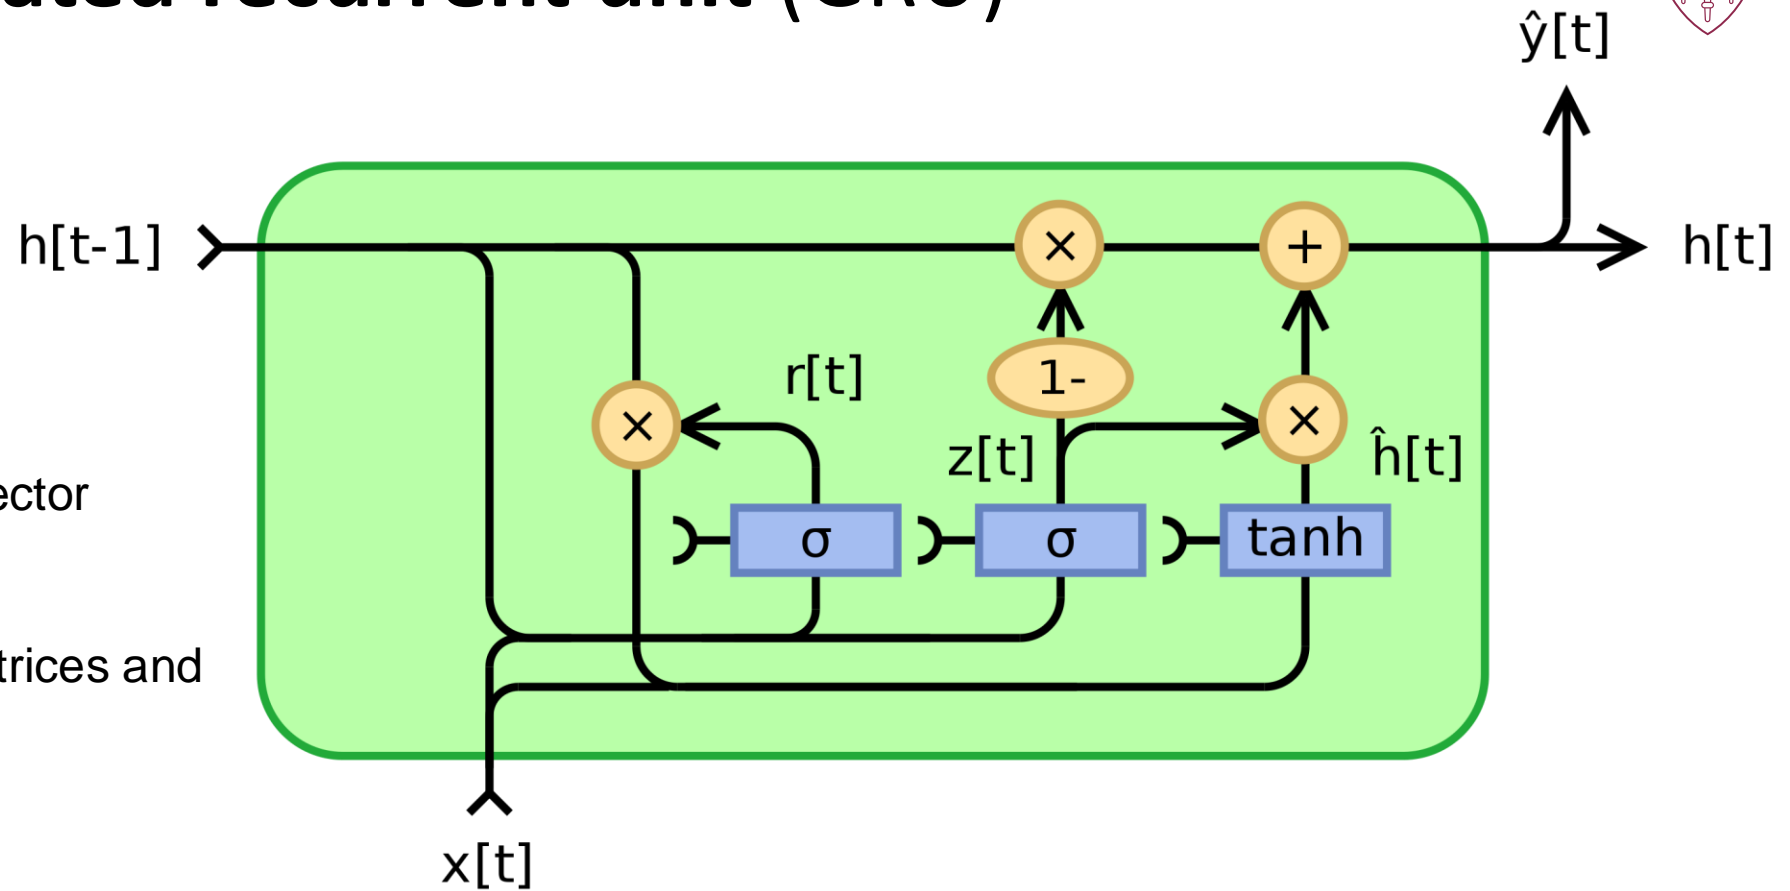
# Appendix



# Gated recurrent unit (GRU)



- $x_t$ : input vector
- $h_t$ : output vector
- $\hat{h}_t$ : candidate activation vector
- $z_t$ : update gate vector
- $r_t$ : reset gate vector
- $W$ ,  $U$  and  $b$ : parameter matrices and vector



(O-like symbol: element-wise multiplication)

<https://ahmedabadmirror.indiatimes.com/entertainment/hollywood/the-gru-hero/articleshow/64033600.cms>



# GRU: Deep Dive

- What do GRUs do?
  - Models can “forget” some of the past ( $Z_t$ ): directly remove past state
  - Models can “reset” ( $r_t$ ): depend more or less on current features
  - Model outputs a (“hidden”) state  $h_t$ : a sum of current features and past states, this is what we feed to next network layer
- Summary: we train a model to care more or less about past depending on current state (provides **importance** of past states)
- This is the simplest solution to solve the vanishing gradient problem
- This method works well, especially for small data
- For larger datasets, LSTM is recommended